

Kevin Takano

SENIOR AI/ML ENGINEER

Curitiba, Parana, Brazil | kevin.takano@gmail.com | linkedin.com/in/ktakanopy | substack.com/@ktakanopy

SUMMARY

Senior engineer with 10 years of software engineering experience, including 5 years building machine learning systems across agentic AI, large language model (LLM) applications, recommender systems, forecasting, and MLOps. Built evaluation loops, guardrails, reusable platforms, and autonomous end-to-end workflows across Scala, Python, TypeScript, Java, GCP, AWS, and distributed systems.

CORE SKILLS

Agentic AI	AI agents, agentic workflows, LLM evaluation, guardrails, feedback loops, OpenAI, Claude Skills, Hamming, LangGraph, LangChain, DSPy, prompt engineering, RAG
Engineering Platforms	Scala, Python, TypeScript, Java, Node.js, SQL, backend systems, distributed systems, APIs, unit testing, CI/CD
Execution	Spark, GCP, BigQuery, AWS, SageMaker, Airflow, Kubernetes, Docker, Postgres, AWS SQS, model monitoring, technical documentation
Execution	Autonomous delivery, end-to-end feature development, product thinking, product ownership, stakeholder communication, async documentation, experimentation

EXPERIENCE

AI Engineer | Luma Health May 2026 - Present

- Developed automated evaluation tests and analysis workflows with Claude Skills and Hamming for a healthcare scheduling agent, identifying critical bugs and improving quality before production release.
- Designed, developed, and deployed a new FAQ agent architecture that improved response quality by 10-15% and made the agent approximately 3x faster.

Senior Machine Learning Engineer | QuintoAndar May 2025 - May 2026

- Built a production LLM evaluation workflow for rent inspections using images and policy rules, with metrics, alerts, database writes, and monitoring.
- Created an LLM-as-judge feedback loop that improved classification quality by 10% and accelerated prompt iteration through stronger evaluation guardrails and production reliability checks.
- Analyzed internal AI-agent architecture and built multi-step conversation scoring pipelines for agentic workflows, improving quality score from 0.48 to 0.90 using LangGraph, Langfuse, DSPy, Postgres, and AWS SQS.

Machine Learning Engineer | Data Scientist | Nubank May 2024 - May 2025

- Built a Scala and Spark calibration model for credit-card interest-rate prediction, reducing incorrect predictions by 80%.
- Improved engineering velocity with reusable libraries, unit tests, self-serve async docs and specs, faster Docker-based CI/CD, and lower-GPU TensorFlow workflows.
- Connected model behavior to business outcomes through time-series, causal, and revenue metrics with automated reporting pipelines.

Machine Learning Engineer | Data Scientist | Hotmart Apr 2022 - May 2024

- Led an end-to-end embedding recommender system on SageMaker and Kubernetes, increasing click conversion to 8.6% versus a 2.07% baseline and driving 3.8x more purchases.
- Built a treasury forecasting system for PIX and billet payments that reduced error by 5% with a simpler, lower-maintenance feature set.
- Re-architected chatbot context and document-processing flows with Python, RAG, LLMs, Java, Spring Boot, and LangChain, while maintaining shared MLOps tooling.

Data Scientist Freelancer | Independent Oct 2024 - Jan 2025

- Built TypeScript and Node.js LLM agents for Twitter and Telegram with plugin development, external tool actions, and OpenAI or OpenRouter integrations.
- Delivered an end-to-end CatBoost conversion model with feature engineering, optimization, and deployment pipeline design that reduced costs by 50%.

EARLIER EXPERIENCE

Data Scientist | Instituto de Pesquisas Eldorado 2020 - 2022

Built dashboards, supervised and unsupervised models, and BigQuery pipelines for digital-bank acquisition; clustering reached 5% conversion versus 1% from random targeting.

Data Engineer | Linx 2018 - 2020

Built Scala and Spark pipelines for segmentation, identity resolution, and recommendations; reduced AWS ingestion jobs from hours to minutes with incremental processing.

Back-End Software Engineer | SIDIA - Samsung Instituto de P&D da Amazonia 2016 - 2018

Developed Java, C++, and Node.js components for Android and Ginga middleware, plus CI and test automation that cut test runtime from hours to seconds and reduced library size by 50%.

PORTFOLIO PROJECTS

CAPY Appointment Bot | github.com/ktakanopy/capy-appointment-bot

Built a clinic appointment assistant with FastAPI, LangGraph, OpenAI, and Streamlit for secure, multi-turn identity verification and appointment management.

- Kept authorization and mutations deterministic with verification gating, retry lockout, ownership checks, idempotent actions, and PII-redacted Langfuse tracing.
- Created unit, graph, API, and multi-turn eval suites with an LLM judge to validate rerouting, failure recovery, and prompt-injection resistance.

Recommendation AI Portfolio | github.com/ktakanopy/recommendation-ai

Built a recommendation-system portfolio spanning research notebooks, a production FastAPI service, and a LangGraph Telegram bot for onboarding and personalized movie recommendations.

- Implemented a PyTorch neural collaborative filtering study with feature engineering, cold-start handling, and evaluation metrics such as Hit Ratio@10 and MRR.
- Built a production-ready API with Clean Architecture, PostgreSQL persistence, Annoy vector indexing, JWT authentication, and hybrid candidate generation.
- Added a user-facing conversational bot that collects demographics and ratings, then delivers recommendations through a multi-step LangGraph workflow.

EDUCATION

Master's in Informatics, Machine Learning | Universidade Federal do Amazonas

2019 - 2022

Bachelor of Computer Science | Universidade Federal do Amazonas

2012 - 2016